



**CS464**  
**Machine Learning**  
**HW2 Report**

**N. Onur VURAL**

**21902330**

**SEC: 2**

# CS464

## Machine Learning - HW2

N. Onur VURAL - 21902330

SEC: 2

### QUESTION 1: PCA & DOGS

```
In [1]: # import necessary Libraries
import os
from os import listdir
import PIL
from PIL import Image
import numpy as np
import matplotlib.pyplot as plt
```

### Converting dog images to 3-D array of size 5239×4096×3

```
In [2]: dog_images = np.zeros((5239,4096,3))
index = 0
# resizing images
folder_dir = "afhq_dog"
for image in os.listdir(folder_dir):
    if (image.endswith(".jpg")):
        resizedIm = PIL.Image.open("afhq_dog/"+ image).resize((64,64), Image.BILINEAR)
        #plt.imshow(resizedIm)
        resizedIm = np.array(resizedIm)
        # resizedIm = np.array(resizedIm, dtype = "int32")

        # print(resizedIm.shape)
        d1, d2, d3 = resizedIm.shape
        resizedIm = resizedIm.flatten().reshape(d1*d2, d3)
        # print(resizedIm)
        # print(resizedIm.shape)
        dog_images[index] = resizedIm
        index = index + 1

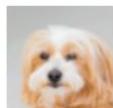
print(dog_images.shape)
```

```
C:\Users\user\AppData\Local\Temp\ipykernel_2508\1957072623.py:7: DeprecationWarning: BILINEAR
is deprecated and will be removed in Pillow 10 (2023-07-01). Use Resampling.BILINEAR instead.
    resizedIm = PIL.Image.open("afhq_dog/"+ image).resize((64,64), Image.BILINEAR)

(5239, 4096, 3)
```

```
In [3]: # sample image
image_id = 0
sample_image = (dog_images[image_id]).astype(np.uint8).reshape(64,64,3)
Image.fromarray(sample_image)
```

Out[3]:



## Q1.1: Applying PCA to get first 10 principal components for each image

```
In [4]: print(dog_images[:, :, 2])
print(dog_images[:, :, 2].shape)
#-----
def PCA(channel_data):

    # PROCESS 1: Zero centering standardaziton
    avg = np.mean(channel_data, axis=0)
    # print("Mean vector is: ", avg)
    # print("Mean vector size is: ", avg.size)
    channel_data2 = channel_data - avg

    # PROCESS 2: Covariance matrix
    covariance_matrix = np.cov(channel_data2, rowvar = False)
    # print(covariance_matrix.shape)
    eig_values, eig_vectors = np.linalg.eig(covariance_matrix)
    # we need eigens in sorted decreasing order
    sorted_order = np.argsort(eig_values)[::-1]
    sorted_eig_values = eig_values[sorted_order]
    sorted_eig_vectors = eig_vectors[:,sorted_order]

    return sorted_eig_values, sorted_eig_vectors # eig_vectors_pc, avg
```

```
[[184. 184. 184. ... 71. 77. 88.]
 [216. 222. 192. ... 80. 71. 70.]
 [123. 124. 123. ... 44. 42. 36.]
 ...
 [100. 100. 96. ... 141. 141. 116.]
 [ 15.  16.  21. ... 201. 199. 187.]
 [188. 178. 172. ... 165. 175. 169.]]
(5239, 4096)
```

```
In [ ]:
```

## Testing for R Channel

```
In [5]: pca_components_num = 10
sorted_eig_values_r, sorted_eig_vectors_r = PCA(dog_images[:, :, 0])

# bring k of them as principle components
eig_vectors_pc_r = sorted_eig_vectors_r[:, 0:pca_components_num]
eig_values_pc_r = sorted_eig_values_r[0:pca_components_num]

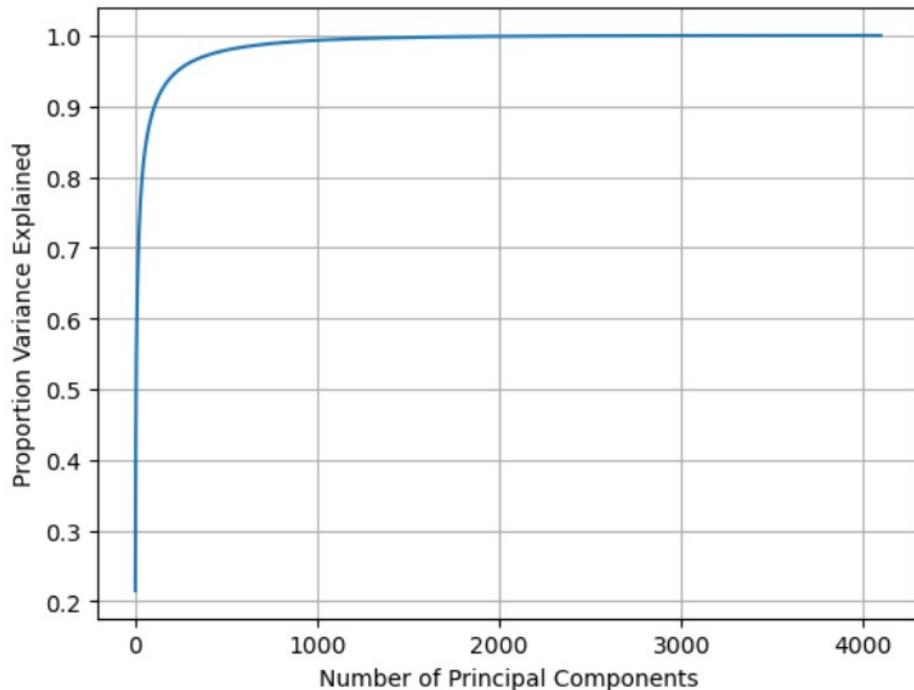
pve = []
for i in range(len(sorted_eig_values_r)):
    pve.append(sorted_eig_values_r[i] / np.sum(sorted_eig_values_r))
# Reporting PVE's, PVE sum, all PVE sum
print("First ", pca_components_num, " pve's are: ", pve[0:pca_components_num])
print("Sum of first", pca_components_num, " pve's is: ", np.sum(pve[0:pca_components_num]))
print("Sum of ALL pve's is: ", np.sum(pve))

# Finding minimum number of PC's to reach above 70% PVE
sum_so_far = 0
i = 0
while sum_so_far < 0.7:
    sum_so_far = sum_so_far + pve[i]
    i = i + 1
print("Minimum number of PC's to obtain 70% PVE is: ", i)

plt.grid()
plt.plot(np.cumsum(pve))
plt.xlabel('Number of Principal Components')
plt.ylabel('Proportion Variance Explained')
```

```
First 10 pve's are: [0.2150681068731377, 0.13542095903812393, 0.0750408239301227, 0.05172646440945615, 0.04228590064660128, 0.024580425792267956, 0.021770322002457612, 0.019896569510316386, 0.01706974941947941, 0.016558493563630355]
Sum of first 10 pve's is: 0.6194178151855935
Sum of ALL pve's is: 0.9999999999999999
Minimum number of PC's to obtain 70% PVE is: 18
```

Out[5]: Text(0, 0.5, 'Proportion Variance Explained')



In [ ]:

## Testing for G Channel

```
In [6]: sorted_eig_values_g, sorted_eig_vectors_g = PCA(dog_images[:, :, 1])

# bring k of them as principle components
eig_vectors_pc_g = sorted_eig_vectors_g[:, 0:pca_components_num]
eig_values_pc_g = sorted_eig_values_g[0:pca_components_num]

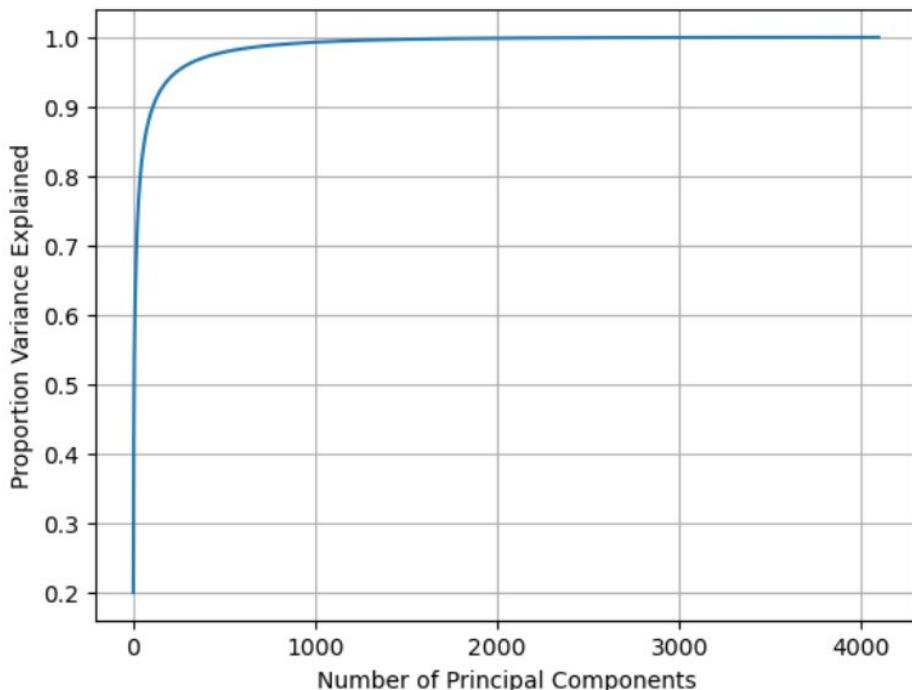
pve = []
for i in range(len(sorted_eig_values_g)):
    pve.append(sorted_eig_values_g[i] / np.sum(sorted_eig_values_g))
# Reporting PVE's, PVE sum, all PVE sum
print("First ", pca_components_num, " pve's are: ", pve[0:pca_components_num])
print("Sum of first", pca_components_num, " pve's is: ", np.sum(pve[0:pca_components_num]))
print("Sum of ALL pve's is: ", np.sum(pve))

# Finding minimum number of PC's to reach above 70% PVE
sum_so_far = 0
i = 0
while sum_so_far < 0.7 :
    sum_so_far = sum_so_far + pve[i]
    i = i + 1
print("Minimum number of PC's to obtain 70% PVE is: ", i)

plt.grid()
plt.plot(np.cumsum(pve))
plt.xlabel('Number of Principal Components')
plt.ylabel('Proportion Variance Explained')
```

```
First 10 pve's are: [0.20045373100504224, 0.13767588143152348, 0.07695187885576901, 0.053969649951491054, 0.042918143529197046, 0.02602155770284145, 0.02142609389411923, 0.020812494401399997, 0.017393203864107407, 0.016811114431734317]
Sum of first 10 pve's is: 0.6144337490672251
Sum of ALL pve's is: 0.9999999999999998
Minimum number of PC's to obtain 70% PVE is: 19
```

```
Out[6]: Text(0, 0.5, 'Proportion Variance Explained')
```



```
In [ ]:
```

## Testing for B Channel

```
In [7]: sorted_eig_values_b, sorted_eig_vectors_b = PCA(dog_images[:, :, 2])

# bring k of them as principle components
eig_vectors_pc_b = sorted_eig_vectors_b[:, 0:pca_components_num]
eig_values_pc_b = sorted_eig_values_b[0:pca_components_num]

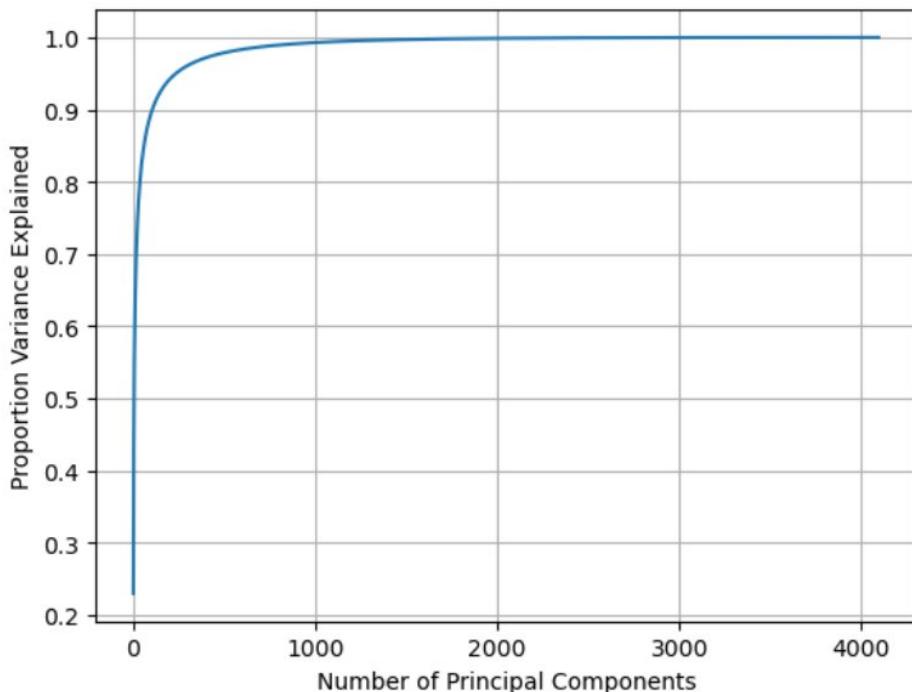
pve = []
for i in range(len(sorted_eig_values_b)):
    pve.append(sorted_eig_values_b[i] / np.sum(sorted_eig_values_b))
# Reporting PVE's, PVE sum, all PVE sum
print("First ", pca_components_num, " pve's are: ", pve[0:pca_components_num])
print("Sum of first", pca_components_num, " pve's is: ", np.sum(pve[0:pca_components_num]))
print("Sum of ALL pve's is: ", np.sum(pve))

# Finding minimum number of PC's to reach above 70% PVE
sum_so_far = 0
i = 0
while sum_so_far < 0.7:
    sum_so_far = sum_so_far + pve[i]
    i = i + 1
print("Minimum number of PC's to obtain 70% PVE is: ", i)

plt.grid()
plt.plot(np.cumsum(pve))
plt.xlabel('Number of Principal Components')
plt.ylabel('Proportion Variance Explained')
```

First 10 pve's are: [0.2299456213159542, 0.13677010354278865, 0.0703323247123188, 0.05355895002957555, 0.03981729858197309, 0.02373057978965713, 0.02098964613416299, 0.02075654086553708, 0.016680265709240864, 0.016291470910427007]  
Sum of first 10 pve's is: 0.6288728015916353  
Sum of ALL pve's is: 1.0000000000000002  
Minimum number of PC's to obtain 70% PVE is: 17

Out[7]: Text(0, 0.5, 'Proportion Variance Explained')



In [ ]:

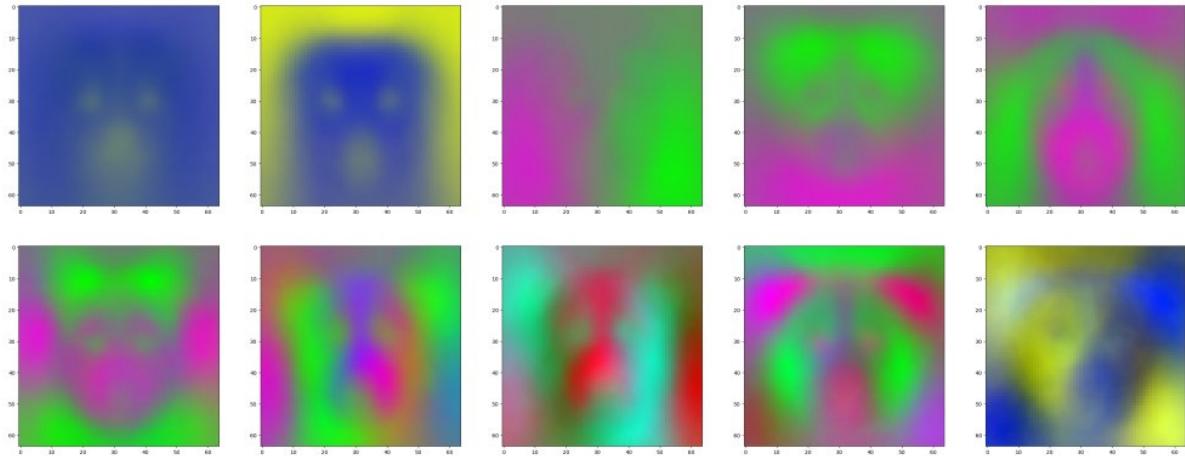
## Q1.2: Obtaining RGB images from 10 Principal Components

```
In [8]: # Reshaping the 10 principal components into 64*64 form
PCMatrices_r = eig_vectors_pc_r.T.reshape(10,64,64)
PCMatrices_g = eig_vectors_pc_g.T.reshape(10,64,64)
PCMatrices_b = eig_vectors_pc_b.T.reshape(10,64,64)

# Normalize values between 0 and 1 by Min-Max Scaling
PCMatrices_r = (PCMatrices_r - PCMatrices_r.min()) / (PCMatrices_r.max() - PCMatrices_r.min())
PCMatrices_g = (PCMatrices_g - PCMatrices_g.min()) / (PCMatrices_g.max() - PCMatrices_g.min())
PCMatrices_b = (PCMatrices_b - PCMatrices_b.min()) / (PCMatrices_b.max() - PCMatrices_b.min())
print(PCMatrices_r.shape)

# obtain 10 RGB images of size 64 x 64 x 3
fig = plt.figure(figsize=[40,40])
for i in range(len(PCMatrices_r)):
    rgb_image = np.dstack((PCMatrices_r[i],PCMatrices_g[i],PCMatrices_b[i]))
    fig.add_subplot(5,5,i+1)
    plt.imshow(rgb_image)
plt.show()
```

(10, 64, 64)



The results demonstrate that those 10 principal components correspond to eigen vector images. Eigen vector images are essentially sufficient to represent all images within the dataset with a specific linear combination.

```
In [ ]:
```

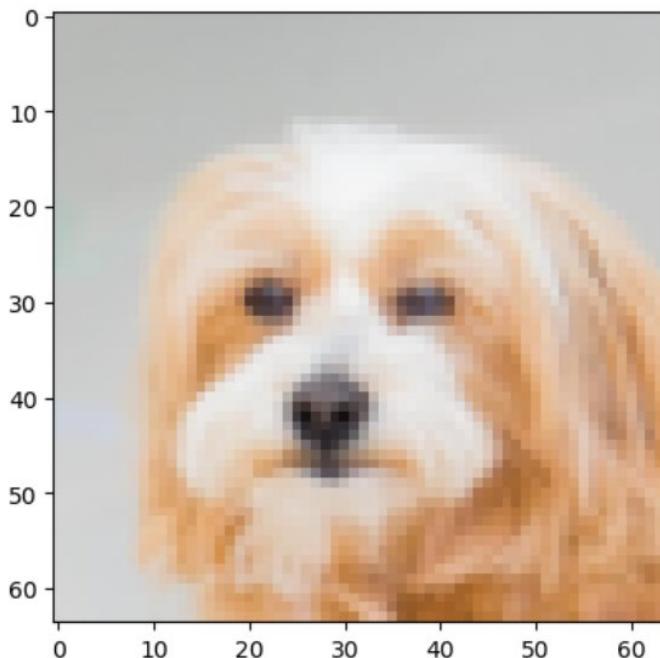
## Q1.3 Reconstructing Original Images with Various Principal Component Values

To reconstruct an original image by using the principal components,

```
In [9]: # Reconstruct first image!!!
first_image_compressed = dog_images[0]
print(first_image_compressed)
plt.imshow(first_image_compressed.reshape(64,64,3).astype("uint8"))
```

```
[[186. 186. 184.]
 [186. 186. 184.]
 [185. 187. 184.]
 ...
 [190. 122. 71.]
 [194. 126. 77.]
 [202. 134. 88.]]
```

```
Out[9]: <matplotlib.image.AxesImage at 0x179474a4ca0>
```



```
In [ ]:
```

## FORMAL RECONSTRUCTION PROCESS

```
In [10]: # FORMAL RECONSTRUCTION PROCESS WITH k set!
# split the image into channels
red_channel = first_image_compressed[:, 0]
green_channel = first_image_compressed[:, 1]
blue_channel = first_image_compressed[:, 2]

# define our k set where each k corresponds to number of pc's
k_set = (1, 50, 250, 500, 1000, 4096)

for k in k_set:
    print("FOR K=", k)
    # Bring corresponding k principal components among all PC's
    # bring k of them as principle components
    eig_vectors_pc_r = sorted_eig_vectors_r[:, 0:k]
    eig_vectors_pc_g = sorted_eig_vectors_g[:, 0:k]
    eig_vectors_pc_b = sorted_eig_vectors_b[:, 0:k]

    # Reshaping the k principal components into 64*64 form
    PCMatrices_r = eig_vectors_pc_r.T.reshape(k, 64, 64)
    PCMatrices_g = eig_vectors_pc_g.T.reshape(k, 64, 64)
    PCMatrices_b = eig_vectors_pc_b.T.reshape(k, 64, 64)

    # TAKING DOT PRODUCT FOR ALL 3 CHANNELS
    dot_product_pc_im_r = np.dot(PCMatrices_r.reshape(k, 4096), red_channel)
    dot_product_pc_im_r = np.matmul(PCMatrices_r.reshape(k, 4096).T, dot_product_pc_im_r)

    # for G
    dot_product_pc_im_g = np.dot(PCMatrices_g.reshape(k, 4096), green_channel)
    dot_product_pc_im_g = np.matmul(PCMatrices_g.reshape(k, 4096).T, dot_product_pc_im_g)

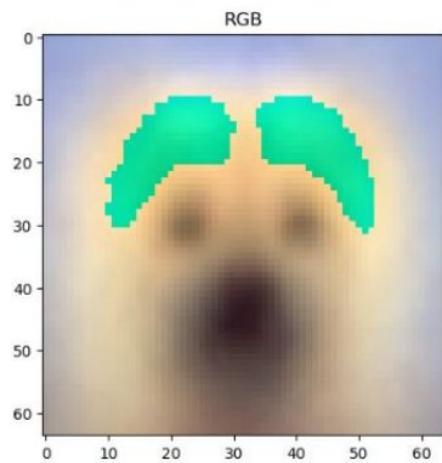
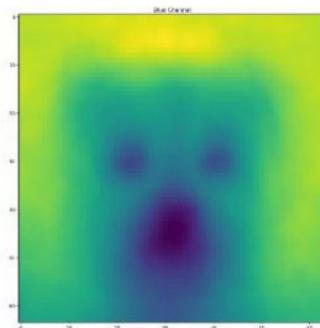
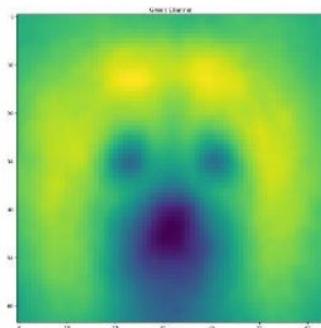
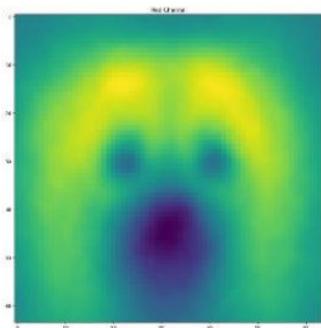
    # for B
    dot_product_pc_im_b = np.dot(PCMatrices_b.reshape(k, 4096), blue_channel)
    dot_product_pc_im_b = np.matmul(PCMatrices_b.reshape(k, 4096).T, dot_product_pc_im_b)

    # OBTAIN RGB IMAGE BY COMBINING ALL PROCESSED CHANNELS
    my_rgb_image = np.dstack((dot_product_pc_im_r, dot_product_pc_im_g, dot_product_pc_im_b))

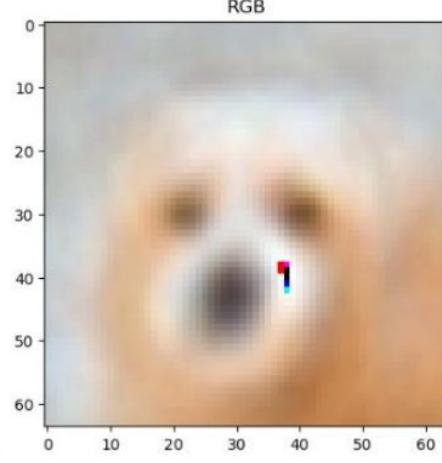
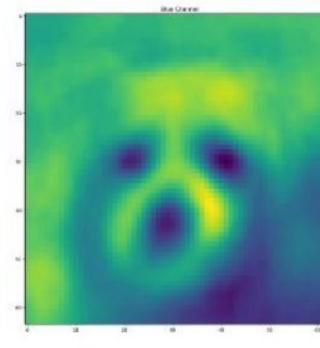
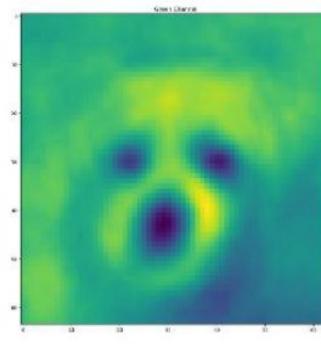
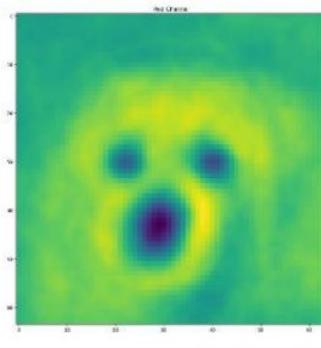
    # PRINTING THE RESULTS
    fig = plt.figure(figsize=[40, 40])
    fig.add_subplot(3, 3, 1)
    plt.title("Red Channel")
    plt.imshow(dot_product_pc_im_r.reshape(64, 64))
    fig.add_subplot(3, 3, 2)
    plt.title("Green Channel")
    plt.imshow(dot_product_pc_im_g.reshape(64, 64))
    fig.add_subplot(3, 3, 3)
    plt.title("Blue Channel")
    plt.imshow(dot_product_pc_im_b.reshape(64, 64))
    plt.show()

    plt.title("RGB")
    plt.imshow(my_rgb_image.reshape(64, 64, 3).astype("uint8"))
    plt.show()
```

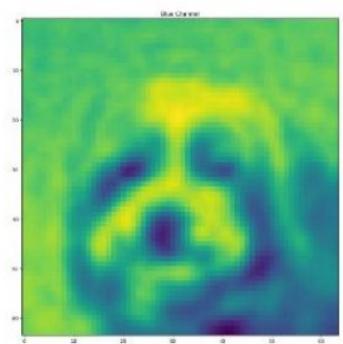
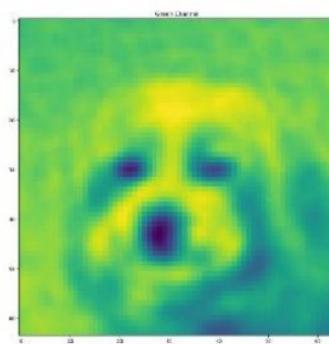
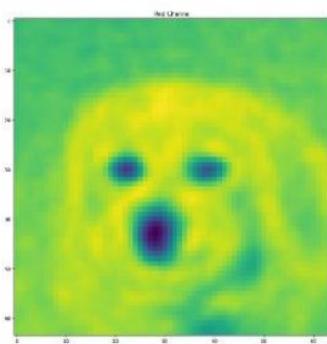
FOR K= 1



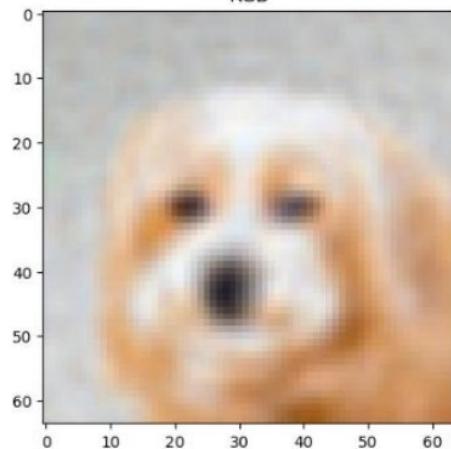
FOR K= 50



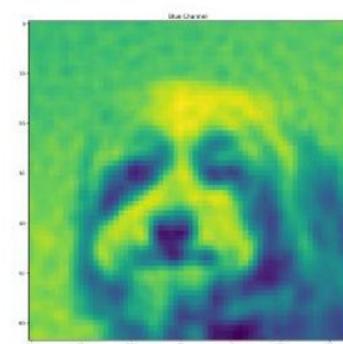
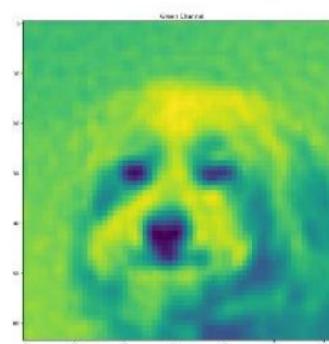
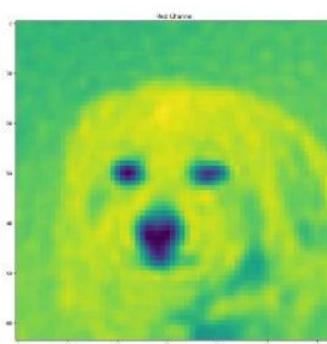
FOR K= 250



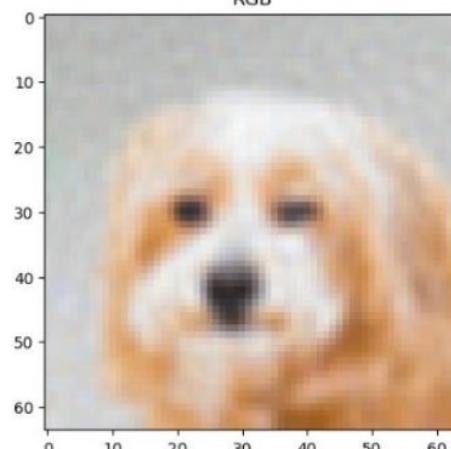
RGB



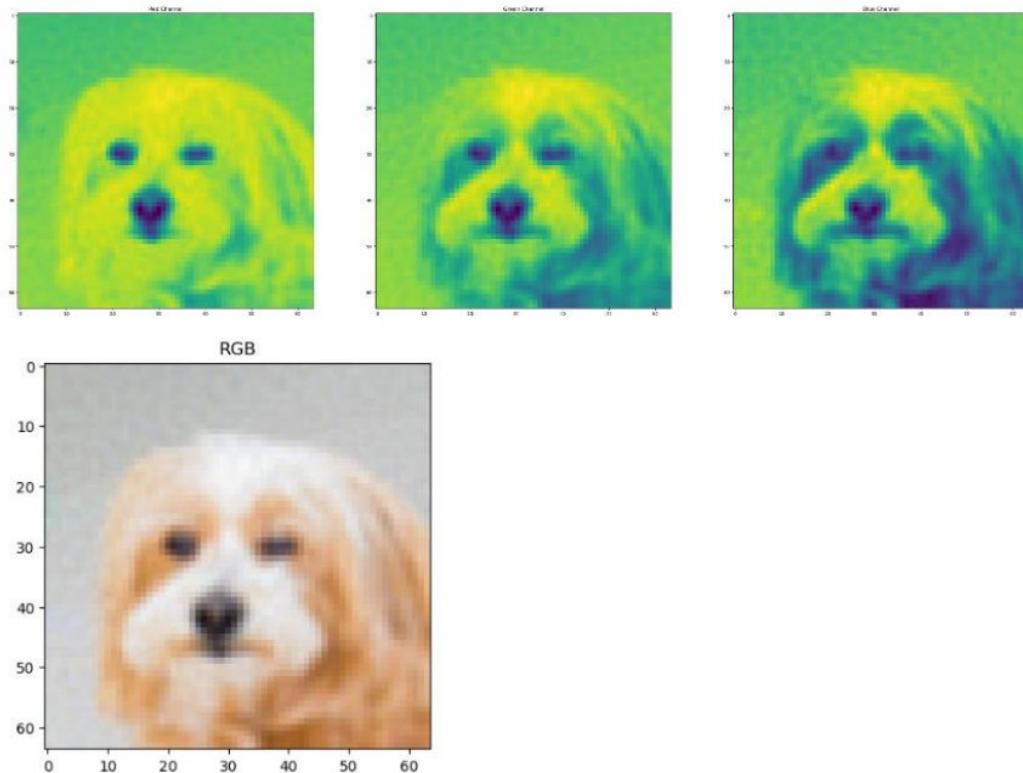
FOR K= 500



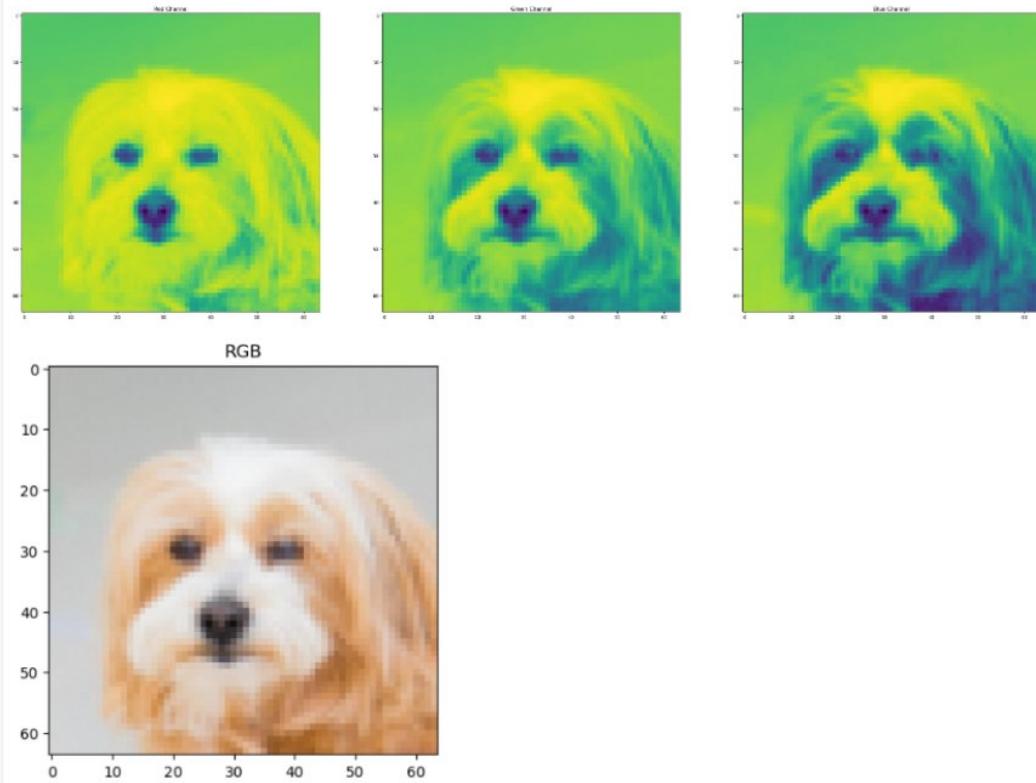
RGB



FOR K= 1000



FOR K= 4096



As the results demonstrate, increasing the number of principal components contribute to a more true representation of the given image. This was expected since the original image is obtained by using linear combination of our principal components. The higher number of principal components increases the accuracy rate of reconstruction.

## QUESTION 2 - Logistic Regression

### STEP 1: OBTAINING DATA

```
In [1]: import pandas as pd
import numpy as np
import math

grid_df = pd.read_csv("dataset.csv")
print(grid_df)

grid_df_arr = np.array(grid_df)
print(grid_df_arr)
print(grid_df_arr.shape)

    tau1      tau2      tau3      tau4      p1      p2      p3  \
0    2.959060  3.079885  8.381025  9.780754  3.763085 -0.782604 -1.257395
1    9.304097  4.902524  3.047541  1.369357  5.067812 -1.940058 -1.872742
2    8.971707  8.848428  3.046479  1.214518  3.405158 -1.207456 -1.277210
3    0.716415  7.669600  4.486641  2.348563  3.963791 -1.027473 -1.938944
4    3.134112  7.608772  4.943759  9.857573  3.525811 -1.125531 -1.845975
...     ...      ...      ...      ...      ...      ...      ...
59995  2.930406  2.376523  9.487627  6.187797  3.343416 -1.449106 -0.658054
59996  3.392299  2.954947  1.274827  6.894759  4.349512 -0.952437 -1.663661
59997  2.364034  8.776391  2.842030  1.008906  4.299976 -0.943884 -1.380719
59998  9.631511  2.757071  3.994398  7.821347  2.514755 -0.649915 -0.966330
59999  6.530527  4.349695  6.781790  8.673138  3.492807 -1.532193 -1.390285

        p4      g1      g2      g3      g4  label
0    -1.723086  0.650456  0.859578  0.887445  0.958034  0
1    -1.255012  0.413441  0.862414  0.562139  0.781760  1
2    -0.920492  0.163041  0.766689  0.839444  0.109853  0
3    -0.997374  0.446209  0.976744  0.929381  0.362718  0
4    -0.554305  0.797110  0.455450  0.656947  0.820923  0
...     ...      ...      ...      ...
59995 -1.236256  0.601709  0.813512  0.779642  0.608385  0
59996 -1.733414  0.502079  0.285880  0.567242  0.366120  1
59997 -1.975373  0.487838  0.149286  0.986505  0.145984  1
59998 -0.898510  0.365246  0.889118  0.587558  0.818391  0
59999 -0.570329  0.073056  0.378761  0.505441  0.942631  0

[60000 rows x 13 columns]
[[2.95906002 3.0798852  8.38102539 ... 0.88744492 0.95803399 0.
[9.30409723 4.90252411 3.04754073 ... 0.56213905 0.78175991 1.
[8.97170691 8.84842842 3.04647875 ... 0.83944402 0.10985324 0.
...
[2.36403419 8.77639096 2.84203025 ... 0.98650532 0.14598403 1.
[9.63151069 2.75707093 3.9943976 ... 0.58755755 0.81839133 0.
[6.53052662 4.34969522 6.7817899 ... 0.50544105 0.94263083 0.
(60000, 13)
```

```
In [2]: def splitData(data, train_size_ratio, val_size_ratio, test_size_ratio):
    data_arr = np.array(data)

    number_of_instances = len(data_arr)

    data_arr_x = data_arr[:, :-1]
    data_arr_y = data_arr[:, -1]

    train_size = int(number_of_instances * train_size_ratio)
    val_size = int(number_of_instances * val_size_ratio)
    test_size = int(number_of_instances * test_size_ratio)

    train_x = data_arr_x[:train_size]
    train_y = data_arr_y[:train_size]

    val_x = data_arr_x[train_size: train_size + val_size]
    val_y = data_arr_y[train_size: train_size + val_size]

    test_x = data_arr_x[train_size + val_size:]
    test_y = data_arr_y[train_size + val_size:]

    print(len(train_x), len(train_y) )
    print(len(val_x), len(val_y) )
    print(len(test_x), len(test_y) )

    return train_x, train_y.astype(int), val_x, val_y.astype(int), test_x, test_y.astype(int)
```

```
In [3]: def normalizeData(data_arr, min_val, max_val):
    normalized_data = (data_arr - min_val) / (max_val - min_val)
    return normalized_data
```

```
In [4]: train_x, train_y, val_x, val_y, test_x, test_y = splitData(grid_df, 0.7, 0.1, 0.2)

max_val = train_x.max()
min_val = train_x.min()

# WE NEED TO NORMALIZE DATA TO PREVENT FEATURE SCALE DIFFERENCE EFFECT
train_x = normalizeData(train_x, min_val, max_val)
val_x = normalizeData(val_x, min_val, max_val)
test_x = normalizeData(test_x, min_val, max_val)

42000 42000
6000 6000
12000 12000
```

## Implementing the Logistic Regression Classifier

```
In [5]: # data_x = (num, 12)
# w = w1, w2, w3, w4....
# bias = w0
# return 0 or 1
class LogisticRegressionClassifier:
    def __init__(self, learning_rate, batch_size):
        self.learning_rate = learning_rate
        self.batch_size = batch_size
        #self.number_epoch = number_epoch
        self.w = None
        self.bias = None

    def MCLE(self, data_x, w, bias):
        sumTerm = np.dot(data_x, w) + bias
        return 1 / (1 + np.exp(sumTerm))

    def init_w_Zeros(self, num):
        x = np.zeros(num)
        self.w, self.bias = x[1:], x[0]

    def init_w_Gaussian(self, mean, standard_dev, num):
        x = np.random.normal(loc = mean, scale = standard_dev, size = num )
        self.w, self.bias = x[1:], x[0]

    def init_w_Uniform(self, num):
        x = np.random.uniform(size = num)
        self.w, self.bias = x[1:], x[0]

    def fit(self, train_x, train_y):
        # implementing the UPDATE RULE for any given batch size
        number_samples, number_features = train_x.shape
        i = 0
        while i < number_samples:
            train_x_slice = train_x[i: self.batch_size + i] # one row = one sample
            train_y_slice = train_y[i: self.batch_size + i]
            i = i + self.batch_size

            y_predicted = self.MCLE(train_x_slice, self.w, self.bias)
            y_predicted_label1 = 1 - y_predicted
            y_real_y_predicted_diff = train_y_slice - y_predicted_label1

            gradient = np.dot(train_x_slice.T, y_real_y_predicted_diff)
            self.w = self.w + self.learning_rate * gradient

            g = np.sum(y_real_y_predicted_diff)
            self.bias = self.bias + self.learning_rate * g

        def predict(self, test_x):
            model_results = self.MCLE(test_x, self.w, self.bias)
            y_predicted = [0 if i > 0.5 else 1 for i in model_results]
            return y_predicted

        def confusionMatrix(self, data_x, data_y):
            predictions = self.predict(data_x)
            labelVarNum = np.unique(data_y).size # number of class variations
            confusionMatrix = np.zeros(shape=(labelVarNum, labelVarNum))
            for pred, exp in zip(predictions, data_y):
                confusionMatrix[pred][exp] += 1
            return confusionMatrix
```

```
In [6]: def accuracyResults(validLabels, validPredictedLabels):
    numTrue = 0
    for gtl, predL in zip(validLabels, validPredictedLabels):
        if gtl == predL:
            numTrue += 1
    numFalse = validLabels.size - numTrue
    #print("The number of correct predictions are", numTrue)
    #print("The number of incorrect predictions are", numFalse)
    accuracy = (numTrue/len(validLabels))
    return accuracy
```



```
In [9]: # full-batch gradient ascent (batch size= number_samples)
learning_rate = 0.001
batch_size = number_samples

LRC3 = LogisticRegressionClassifier(learning_rate, batch_size)
LRC3.init_w_Gaussian(0, 1, total_w)

accuracy_array_batch_full = []
# Make experiment within 100 epoch
for i in range(100):
    # full-batch gradient ascent (batch size= number_samples)
    LRC3.fit( train_x, train_y )
    pred = LRC3.predict(val_x)
    accuracy_array_batch_full.append(accuracyResults(val_y, pred))

print(accuracy_array_batch_full)
```

[0.641, 0.359, 0.641, 0.641, 0.359, 0.641, 0.644833333333334, 0.359, 0.641, 0.649333333333333, 0.359, 0.641, 0.657, 0.359, 0.641, 0.6571666666666667, 0.359, 0.641, 0.660833333333334, 0.359, 0.641, 0.6615, 0.359, 0.641, 0.6625, 0.359, 0.641, 0.6628333333333334, 0.359, 0.641, 0.663, 0.359, 0.641, 0.664333333333333, 0.3591666666666667, 0.641, 0.665333333333333, 0.3616666666666667, 0.641, 0.667333333333333, 0.3705, 0.641, 0.672, 0.3835, 0.641, 0.6845, 0.404833333333333, 0.641, 0.694833333333333, 0.428833333333334, 0.641, 0.7115, 0.476833333333333, 0.641, 0.7286666666666667, 0.61466666666666667, 0.6501666666666667, 0.5356666666666666, 0.6421666666666667, 0.710333333333334, 0.7206666666666667, 0.5098333333333334, 0.641 5, 0.732333333333333, 0.723, 0.713133333333333, 0.6531666666666667, 0.6615, 0.449833333333333, 0.641, 0.7253333333333334, 0.5 2816666666666667, 0.641833333333334, 0.721, 0.730833333333333, 0.6265, 0.653333333333333, 0.507, 0.641333333333333, 0.733, 0. 714333333333333, 0.7226666666666667, 0.503333333333333, 0.6411666666666667, 0.733833333333333, 0.6976666666666667, 0.694833333333333, 0.4311666666666664, 0.641, 0.715333333333334, 0.486, 0.641, 0.7335, 0.648833333333334, 0.6616666666666666, 0.470666666667, 0.641]

According to experiment results, I observed that batch size affected training time as I checked from Jupyter Lab.

More frequent updates mean more expensive computation, more slow to execute

Therefore stochastic is the slowest among all, following mini batch and full batch in terms of exec. time

As in an example run:

- \*Experimenting with batch size = 1 took 69s to execute on my machine
- \*Experimenting with batch size = 64 took 3s to execute on my machine
- \*Experimenting with batch size = 42000 took less th. 1s to execute on my machine

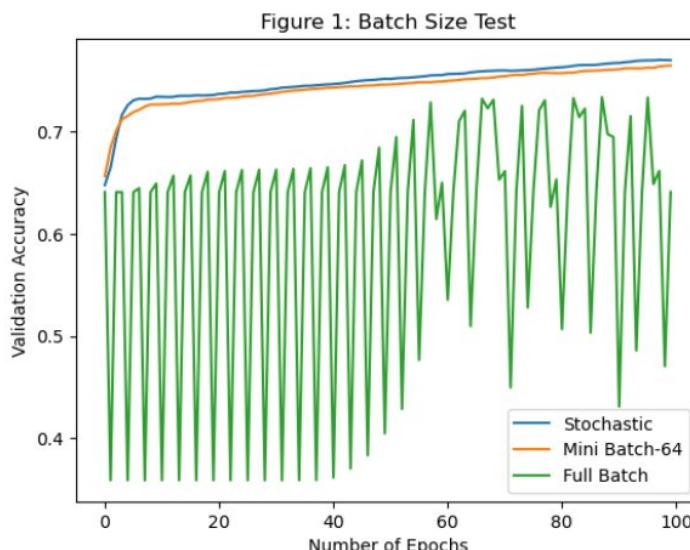
## Plotting graph for batch size experiment results

```
In [10]: import matplotlib.pyplot as plt

x_axis = []
for i in range(100):
    x_axis.append(i)

plt.plot(x_axis, accuracy_array_stochastic, label = "Stochastic")
plt.plot(x_axis, accuracy_array_batch_64, label = "Mini Batch-64")
plt.plot(x_axis, accuracy_array_batch_full, label = "Full Batch")

plt.xlabel('Number of Epochs')
plt.ylabel('Validation Accuracy')
plt.title('Figure 1: Batch Size Test')
plt.legend()
plt.show()
```



At the end of 100 epochs, the best validation accuracy belongs to stochastic & mini batch technique! (very close)  
Here is the confusion matrix of it below:



```
In [14]: # Zero init.
learning_rate = 0.001
batch_size = 64

LRC_zero = LogisticRegressionClassifier(learning_rate, batch_size= 64)
LRC_zero.init_w_Zeros(total_w)

accuracy_array_zero = []
# Make experiment within 100 epoch
for i in range(100):
    LRC_zero.fit(train_x, train_y)
    pred = LRC_zero.predict(val_x)
    accuracy_array_zero.append(accuracyResults(val_y, pred))

print(accuracy_array_zero)

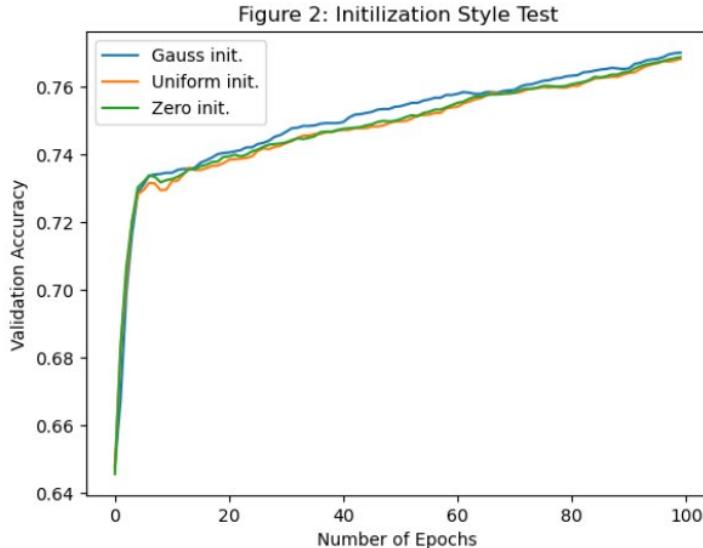
[0.6456666666666667, 0.6831666666666667, 0.7066666666666667, 0.7201666666666666, 0.7303333333333333, 0.732, 0.7338333333333333, 0.7335, 0.7318333333333333, 0.7325, 0.7328333333333333, 0.7335, 0.7345, 0.7356666666666667, 0.7356666666666667, 0.736666666666666667, 0.737, 0.7378333333333333, 0.738, 0.7393333333333333, 0.7395, 0.74, 0.7395, 0.74, 0.7415, 0.7423333333333333, 0.743, 0.7431666666666666, 0.7433333333333333, 0.7436666666666667, 0.7443333333333333, 0.7448333333333333, 0.7446666666666667, 0.745, 0.7455, 0.7466666666666667, 0.7468333333333333, 0.7468333333333333, 0.7473333333333333, 0.7476666666666667, 0.7478333333333333, 0.748, 0.7481666666666666, 0.7486666666666667, 0.749, 0.7496666666666667, 0.75, 0.749833333333334, 0.75, 0.7505, 0.7511666666666666, 0.7516666666666667, 0.7513333333333333, 0.7523333333333333, 0.7526666666666667, 0.7533333333333333, 0.7533333333333333, 0.754, 0.7546666666666667, 0.7553333333333333, 0.7556666666666667, 0.7565, 0.7571666666666667, 0.7571666666666667, 0.7578333333333334, 0.7578333333333334, 0.758, 0.7583333333333333, 0.7583333333333333, 0.7591666666666667, 0.7595, 0.7595, 0.759833333333334, 0.7603333333333333, 0.7603333333333333, 0.7601666666666667, 0.7601666666666667, 0.7603333333333333, 0.760833333333334, 0.7611666666666667, 0.7615, 0.762, 0.763, 0.762833333333334, 0.7631666666666667, 0.7636666666666667, 0.7636666666666667, 0.7643333333333333, 0.7645, 0.7655, 0.7661666666666667, 0.7666666666666667, 0.767, 0.7673333333333333, 0.7676666666666667, 0.7681666666666667, 0.7685, 0.7688333333333334]
```

## Plotting graph for init. techniques experiment results

```
In [15]: x_axis = []
for i in range(100):
    x_axis.append(i)

plt.plot(x_axis, accuracy_array_gauss, label = "Gauss init.")
plt.plot(x_axis, accuracy_array_uniform, label = "Uniform init.")
plt.plot(x_axis, accuracy_array_zero, label = "Zero init.")

plt.xlabel('Number of Epochs')
plt.ylabel('Validation Accuracy')
plt.title('Figure 2: Initialization Style Test')
plt.legend()
plt.show()
```



Apparently, the results suggest that initialization technique did not make a huge difference on the model accuracy. Still, according to the graph, Gauss is most stable

Here is the confusion matrix of it below:

```
In [16]: pred = LRC_gauss.predict(val_x)
print(accuracyResults(val_y, pred))
LRC_gauss.confusionMatrix(val_x, val_y)

0.7701666666666667

Out[16]: array([[3351., 884.],
   [495., 1270.]])
```

## Q 2.3: Experimenting with different learning rate values

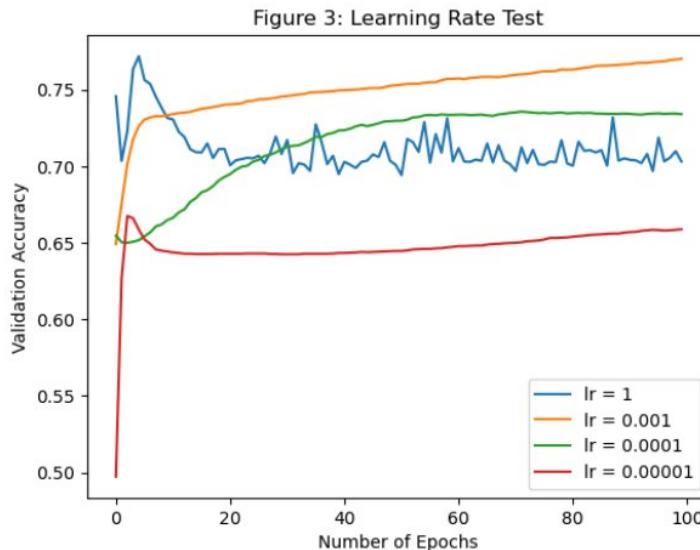




```
In [21]: x_axis = []
for i in range(100):
    x_axis.append(i)

plt.plot(x_axis, accuracy_array_lr_1, label = "lr = 1")
plt.plot(x_axis, accuracy_array_lr_10expminus3, label = "lr = 0.001")
plt.plot(x_axis, accuracy_array_lr_10expminus4, label = "lr = 0.0001")
plt.plot(x_axis, accuracy_array_lr_10expminus5, label = "lr = 0.00001")

plt.xlabel('Number of Epochs')
plt.ylabel('Validation Accuracy')
plt.title('Figure 3: Learning Rate Test')
plt.legend()
plt.show()
```



Learning rate is about how rapidly we should make update for our weights.

When learning rate is too large, overshooting on the optimal value may occur

When learning rate is too small, the number of iterations (epochs) needed to get close to optimal value will increase.

In a constant epoch number (100) case as in our example, learning rate should be selected in such a way that:

- \* It must not be too large for overshooting
- \* It must not be very small so that it cannot converge to optimal value in given 100 iteration time

Accordingly, the results show that:

Largest learning rate oscillates and cannot successfully converge to our data results,

smallest learning rate is incapable of getting close to converge during the given epoch time

In this case 0.001 comes out the most optimal selection choice for our data!

Here is the confusion matrix of it below:

```
In [22]: pred = LRC_lr_10expminus3.predict(val_x)
print(accuracyResults(val_y, pred))
LRC_lr_10expminus3.confusionMatrix(val_x, val_y)
```

0.7701666666666667

```
Out[22]: array([[3348.,  881.],
   [ 498., 1273.]])
```

## OPTIMAL MODEL IS CONSTRUCTED WITH:

BATCH SIZE = 64 (Mini batch is selected over stochastic because the results were very close, mini was much faster)

INIT = GAUSS INIT (All were very close actually - Most stable progress on graph)

LEARNING RATE = 0.001 (Definately the best among others!)

## Q 2.4: Building the optimal model

```
In [23]: # CREATING THE REGRESSOR
number_samples, number_features = train_x.shape
total_w = number_features + 1

# Init with: stochastic, uniform dist, lr = 0.001
batch_size = 64
learning_rate = 0.001

LRC_optimal = LogisticRegressionClassifier(learning_rate, batch_size= 64)
LRC_optimal.init_w_Gaussian(0, 1, total_w)

for i in range(100):
    LRC_optimal.fit(train_x, train_y)

pred = LRC_optimal.predict(test_x)

conf_mat = LRC_optimal.confusionMatrix(test_x, test_y)
print("Confusion matrix is: \n", conf_mat )
print("Accuracy is: %", accuracyResults(test_y, pred)*100)

precision = conf_mat[0][0] / (conf_mat[0][0] + conf_mat[0][1])
recall = conf_mat[0][0] / (conf_mat[0][0] + conf_mat[1][0])
f1_score = ((1 + pow(1, 2)) * precision * recall)/((pow(1,2) * precision) + recall)
f2_score = ((1 + pow(2, 2)) * precision * recall)/((pow(2,2) * precision) + recall)
f05_score = ((1 + pow(1, 0.5)) * precision * recall)/((pow(1,0.5) * precision) + recall)
fp_rate = conf_mat[0][1] / (conf_mat[0][1] + conf_mat[1][1])

print("Precision is: %", precision)
print("Recall is: %", recall)
print("F1 Score is: %", f1_score)
print("F2 Score is: %", f2_score)
print("F0.5 Score is: %", f05_score)
print("FP Rate is: %", fp_rate)
```

```
Confusion matrix is:
[[6652. 1841.]
 [ 999. 2508.]]
Accuracy is: % 76.33333333333333
Precision is: % 0.7832332509125162
Recall is: % 0.8694288328323095
F1 Score is: % 0.8240832507433103
F2 Score is: % 0.8507046576463667
F0.5 Score is: % 0.8240832507433103
FP Rate is: % 0.42331570475971486
```

In our case, to evaluate the overall performance of the model in regard to calculated metrics, we can say that:  
 \*First of all accuracy is important, but it may have limitations when class distributions are not balanced!

\*As our task is to detect unstable grids we would desire to have higher recall,  
 the reason to that is important thing being able to classify as many positives as possible to not miss any positive  
 this is a tradeoff of precision, but falsely predicting a sample as unstable is less risky than missing an unstable one

\*F2 score is more important in this regard as Beta > 1 favors recall!

\*Relative importance of detecting positives and negatives are not same, because as mentioned above  
 detecting positive class that is grid being unstable is important and risky whereas we can tolerate misclassifying  
 some stable grids as unstable, false positive.  
 It would increase the cost to check them again of course, but that is far better than the risk of non detected unstable grid!

```
In [24]: from sklearn.linear_model import LogisticRegression
logisticRegr = LogisticRegression()
for i in range(100):
    logisticRegr.fit(train_x, train_y)
predictions = logisticRegr.predict(test_x)
score = logisticRegr.score(test_x, test_y)
print(score*100)
```

80.575

## CLOSE ENOUGH!!!!

In [ ]: